

# HBASE – PIG INTEGRATION

By.Mr.Gopal Krishna  
Sr.Hadoop Technical Architect  
CCA 175-Spark and Hadoop Developer Certified Consultant

## TO WRITE THE DATA FROM PIG TO HBASE

### FROM HBASE SIDE

+++++

#### VERSIONS USED:

PIG – pig-0.15.0

HBASE – hbase-0.98.4-hadoop2

HADOOP-2.6.0

hbase(main):009:0> create 'hbasepig\_emptable','empdetails'

0 row(s) in 0.5520 seconds

=> Hbase::Table - hbasepig\_emptable

hbase(main):010:0> scan 'hbasepig\_emptable'

ROW COLUMN+CELL

0 row(s) in 0.0470 seconds

hbase(main):011:0>

### FROM PIG SIDE

+++++

#### INPUT DATA:

gopalkrishna@ubuntu:~/PRAC/HBASE\$ cat HBase\_Pig\_InputData.log

1000	RAVI	VARMA	9	TECHNOLOGY LEAD
1001	KOUSIK	MISTRA	11	ENTRYLEVEL ARCHITECT
1002	SAIED	FAISAL	21	DELIVERY HEAD
1003	SOURAV	JAIN	18	SENIOUR PROJECT MANAGER
1004	CHRIS	WAYNE	15	TEST MANAGER
1005	ANEES	FLEXWALA	19	CONFIGURATION CONTROLLER
1006	GOPAL	KRISHNA	13	SR.BIGDATA ARCHITECT
1007	RAMYA	SRI	8	BIGDATA LEAD
1008	MUNI	KRISHNA	5	SR.HADOOP DEVELOPER
1009	RAMYA	SRI	7	SR.SPARK DEVELOPER
1010	RAM	CHANDAR	16	PROJECT MANAGER
1011	AKHILA	PRIYA	3	HADOOP DEVELOPER

```
1012 RAHUL DEV 6 HADOOP ADMINISTRATOR
1013 AMIT MISRA 4 BIGDATA TESTER
1014 PUJA YERUNKAR 11 BIGDATA TEST LEAD
```

**gopalkrishna@ubuntu:~/PRAC/HBASE\$**

### PIG SCRIPT

**gopalkrishna@ubuntu:~/PRAC/HBASE\$ cat HBase\_Pig\_Script.pig**

```
REGISTER /home/gopalkrishna/INSTALL/hbase-0.98.4-hadoop2/lib/zookeeper-3.4.6.jar;
REGISTER /home/gopalkrishna/INSTALL/hbase-0.98.4-hadoop2/lib/guava-12.0.1.jar;
REGISTER /home/gopalkrishna/INSTALL/hbase-0.98.4-hadoop2/lib/protobuf-java-2.5.0.jar;
REGISTER /home/gopalkrishna/INSTALL/hbase-0.98.4-hadoop2/lib/hbase-common-0.98.4-hadoop2.jar;
REGISTER /home/gopalkrishna/INSTALL/hbase-0.98.4-hadoop2/lib/hbase-client-0.98.4-hadoop2.jar;
REGISTER /home/gopalkrishna/INSTALL/hbase-0.98.4-hadoop2/lib/hbase-server-0.98.4-hadoop2.jar;
```

```
iniData = LOAD '/HBase_Pig_Data/HBase_Pig_InputData.log' USING PigStorage('\t') as
    (empid:int,firstname:chararray,lastname:chararray,yrsofexp:int,desg:chararray);
STORE iniData INTO 'hbase://hbasepig_emptable' USING org.apache.pig.backend.hadoop.hbase.HBaseStorage(
    'empdetails:firstname,empdetails:lastname,empdetails:yrsofexp,empdetails:desg');
```

### TO RUN PIG SCRIPT:

**gopalkrishna@ubuntu:~/PRAC/HBASE\$ pig HBase\_Pig\_Script.pig**

---

---

## VERIFICATION FROM HBASE SIDE

+++++

**AFTER RUNNING THE PIG SCRIPT => GO TO HBASE SHELL AND AGAIN RUN THE SCAN COMMAND ON THE HBASE TABLE LIKE BELOW**

**hbase(main):011:0> scan 'hbasepig\_emptable'**

ROW	COLUMN+CELL
1000	column=empdetails:desg, timestamp=1505304875777, value=TECHNOLOGY LEAD
1000	column=empdetails:firstname, timestamp=1505304875777, value=RAVI
1000	column=empdetails:lastname, timestamp=1505304875777, value=VARMA
1000	column=empdetails:yrsofexp, timestamp=1505304875777, value=9
1001	column=empdetails:desg, timestamp=1505304875805, value=ENTRYLEVEL ARCHITECT
1001	column=empdetails:firstname, timestamp=1505304875805, value=KOUSIK
1001	column=empdetails:lastname, timestamp=1505304875805, value=MISTRA
1001	column=empdetails:yrsofexp, timestamp=1505304875805, value=11
1002	column=empdetails:desg, timestamp=1505304875806, value=DELIVERY HEAD
1002	column=empdetails:firstname, timestamp=1505304875806, value=SAIED
1002	column=empdetails:lastname, timestamp=1505304875806, value=FAISAL
1002	column=empdetails:yrsofexp, timestamp=1505304875806, value=21
1003	column=empdetails:desg, timestamp=1505304875807, value=SENIOR PROJECT MANAGER
1003	column=empdetails:firstname, timestamp=1505304875807, value=SOURAV
1003	column=empdetails:lastname, timestamp=1505304875807, value=JAIN
1003	column=empdetails:yrsofexp, timestamp=1505304875807, value=18
1004	column=empdetails:desg, timestamp=1505304875808, value=TEST MANAGER
1004	column=empdetails:firstname, timestamp=1505304875808, value=CHRIS
1004	column=empdetails:lastname, timestamp=1505304875808, value=WAYNE
1004	column=empdetails:yrsofexp, timestamp=1505304875808, value=15
1005	column=empdetails:desg, timestamp=1505304875809, value=CONFIGURATION CONTROLLER
1005	column=empdetails:firstname, timestamp=1505304875809, value=ANEES
1005	column=empdetails:lastname, timestamp=1505304875809, value=FLEXWALA
1005	column=empdetails:yrsofexp, timestamp=1505304875809, value=19

```

1006      column=empdetails:desg, timestamp=1505304875809, value=SR.BIGDATA ARCHITECT
1006      column=empdetails:firstname, timestamp=1505304875809, value=GOPAL
1006      column=empdetails:lastname, timestamp=1505304875809, value=KRISHNA
1006      column=empdetails:yrsofexp, timestamp=1505304875809, value=13
1007      column=empdetails:desg, timestamp=1505304875810, value=BIGDATA LEAD
1007      column=empdetails:firstname, timestamp=1505304875810, value=RAMYA
1007      column=empdetails:lastname, timestamp=1505304875810, value=SRI
1007      column=empdetails:yrsofexp, timestamp=1505304875810, value=8
1008      column=empdetails:desg, timestamp=1505304875811, value=SR.HADOOP DEVELOPER
1008      column=empdetails:firstname, timestamp=1505304875811, value=MUNI
1008      column=empdetails:lastname, timestamp=1505304875811, value=KRISHNA
1008      column=empdetails:yrsofexp, timestamp=1505304875811, value=5
1009      column=empdetails:desg, timestamp=1505304875811, value=SR.SPARK DEVELOPER
1009      column=empdetails:firstname, timestamp=1505304875811, value=RAMYA
1009      column=empdetails:lastname, timestamp=1505304875811, value=SRI
1009      column=empdetails:yrsofexp, timestamp=1505304875811, value=7
1010      column=empdetails:desg, timestamp=1505304875812, value=PROJECT MANAGER
1010      column=empdetails:firstname, timestamp=1505304875812, value=RAM
1010      column=empdetails:lastname, timestamp=1505304875812, value=CHANDAR
1010      column=empdetails:yrsofexp, timestamp=1505304875812, value=16
1011      column=empdetails:desg, timestamp=1505304875813, value=HADOOP DEVELOPER
1011      column=empdetails:firstname, timestamp=1505304875813, value=AKHILA
1011      column=empdetails:lastname, timestamp=1505304875813, value=PRIYA
1011      column=empdetails:yrsofexp, timestamp=1505304875813, value=3
1012      column=empdetails:desg, timestamp=1505304875813, value=HADOOP ADMINISTRATOR
1012      column=empdetails:firstname, timestamp=1505304875813, value=RAHUL
1012      column=empdetails:lastname, timestamp=1505304875813, value=DEV
1012      column=empdetails:yrsofexp, timestamp=1505304875813, value=6
1013      column=empdetails:desg, timestamp=1505304875814, value=BIGDATA TESTER
1013      column=empdetails:firstname, timestamp=1505304875814, value=AMIT
1013      column=empdetails:lastname, timestamp=1505304875814, value=MISRA
1013      column=empdetails:yrsofexp, timestamp=1505304875814, value=4
1014      column=empdetails:desg, timestamp=1505304875815, value=BIGDATA TEST LEAD
1014      column=empdetails:firstname, timestamp=1505304875815, value=PUJA
1014      column=empdetails:lastname, timestamp=1505304875815, value=YERUNKAR
1014      column=empdetails:yrsofexp, timestamp=1505304875815, value=11
15 row(s) in 0.6160 seconds

```

## TO READ THE DATA FROM HBASE TO PIG

---

### FROM PIG SIDE

#### PIG SCRIPT TO READ THE DATA FROM HBASE TABLE

**gopalkrishna@ubuntu:~/PRAC/HBASE\$ cat HBase\_Pig\_Script.pig**

```

REGISTER /home/gopalkrishna/INSTALL/hbase-0.98.4-hadoop2/lib/zookeeper-3.4.6.jar;
REGISTER /home/gopalkrishna/INSTALL/hbase-0.98.4-hadoop2/lib/guava-12.0.1.jar;
REGISTER /home/gopalkrishna/INSTALL/hbase-0.98.4-hadoop2/lib/protobuf-java-2.5.0.jar;
REGISTER /home/gopalkrishna/INSTALL/hbase-0.98.4-hadoop2/lib/hbase-common-0.98.4-hadoop2.jar;
REGISTER /home/gopalkrishna/INSTALL/hbase-0.98.4-hadoop2/lib/hbase-client-0.98.4-hadoop2.jar;
REGISTER /home/gopalkrishna/INSTALL/hbase-0.98.4-hadoop2/lib/hbase-server-0.98.4-hadoop2.jar;

--iniData = LOAD '/HBase_Pig_Data/HBase_Pig_InputData.log' USING PigStorage('\t') as
--      (empid:int,firstname:chararray,lastname:chararray,yrsofexp:int,desg:chararray);
--STORE iniData INTO 'hbase://hbasepig_emptable' USING org.apache.pig.backend.hadoop.hbase.HBaseStorage(
-- 'empdetails:firstname,empdetails:lastname,empdetails:yrsofexp,empdetails:desg');
A = LOAD 'hbase://hbasepig_emptable'

```

```
using
org.apache.pig.backend.hadoop.hbase.HBaseStorage('empdetails:firstname,empdetails:lastname,empdetails:yrsof
exp,empdetails:desg', '-loadKey true')
as (rowkey:bytearray, firstname:chararray,lastname:chararray,yrsofexp:int,desg:chararray);
DUMP A;
```

**gopalkrishna@ubuntu:~/PRAC/HBASE\$**

### **TO RUN PIG SCRIPT:**

**gopalkrishna@ubuntu:~/PRAC/HBASE\$ pig HBase\_Pig\_Script.pig**

**AFTER RUNNING THE PIG SCRIPT => VERIFY THE RESULT ON THE CONSOLE , ALL THE RECORDS IN HBASE TABLE(hbasepig\_emptable) WILL BE DISPLAYED.**

Input(s):

Successfully read 15 records (337 bytes) from: "hbase://hbasepig\_emptable"

Output(s):

Successfully stored 15 records (777 bytes) in: "hdfs://localhost:8020/tmp/temp645148187/tmp-1399112497"

Counters:

Total records written : 15

Total bytes written : 777

Spillable Memory Manager spill count : 0

Total bags proactively spilled: 0

Total records proactively spilled: 0

Job DAG:

job\_1505297080421\_0008

2017-09-13 06:27:24,563 [main] INFO org.apache.hadoop.yarn.client.RMPProxy - Connecting to ResourceManager at localhost/127.0.0.1:8032

2017-09-13 06:27:24,597 [main] INFO org.apache.hadoop.mapred.ClientServiceDelegate - Application state is completed. FinalApplicationStatus=SUCCEEDED. Redirecting to job history server

2017-09-13 06:27:24,684 [main] INFO org.apache.hadoop.yarn.client.RMPProxy - Connecting to ResourceManager at localhost/127.0.0.1:8032

2017-09-13 06:27:24,705 [main] INFO org.apache.hadoop.mapred.ClientServiceDelegate - Application state is completed. FinalApplicationStatus=SUCCEEDED. Redirecting to job history server

2017-09-13 06:27:24,809 [main] INFO org.apache.hadoop.yarn.client.RMPProxy - Connecting to ResourceManager at localhost/127.0.0.1:8032

2017-09-13 06:27:24,830 [main] INFO org.apache.hadoop.mapred.ClientServiceDelegate - Application state is completed. FinalApplicationStatus=SUCCEEDED. Redirecting to job history server

2017-09-13 06:27:24,987 [main] INFO

org.apache.pig.backend.hadoop.executionengine.mapReduceLayer.MapReduceLauncher - Success!

2017-09-13 06:27:24,993 [main] INFO org.apache.hadoop.conf.Configuration.deprecation - fs.default.name is deprecated. Instead, use fs.defaultFS

2017-09-13 06:27:25,001 [main] INFO org.apache.pig.data.SchemaTupleBackend - Key [pig.schematuple] was not set... will not generate code.

2017-09-13 06:27:25,095 [main] INFO org.apache.hadoop.mapreduce.lib.input.FileInputFormat - Total input paths to process : 1

2017-09-13 06:27:25,107 [main] INFO org.apache.pig.backend.hadoop.executionengine.util.MapRedUtil - Total input paths to process : 1

(1000,RAVI,VARMA,9,TECHNOLOGY LEAD)

(1001,KOUSIK,MISTRA,11,ENTRYLEVEL ARCHITECT)

(1002,SAIED,FAISAL,21,DELIVERY HEAD)

(1003,SOURAV,JAIN,18,SENIOR PROJECT MANAGER)

(1004,CHRIS,WAYNE,15,TEST MANAGER)

(1005,ANEES,FLEXWALA,19,CONFIGURATION CONTROLLER)

(1006,GOPAL,KRISHNA,13,SR.BIGDATA ARCHITECT)

(1007,RAMYA,SRI,8,BIGDATA LEAD)

(1008,MUNI,KRISHNA,5,SR.HADOOP DEVELOPER)

(1009,RAMYA,SRI,7,SR.SPARK DEVELOPER)

(1010,RAM,CHANDAR,16,PROJECT MANAGER)

(1011,AKHILA,PRIYA,3,HADOOP DEVELOPER)  
(1012,RAHUL,DEV,6,HADOOP ADMINISTRATOR)  
(1013,AMIT,MISRA,4,BIGDATA TESTER)  
(1014,PUJA,YERUNKAR,11,BIGDATA TEST LEAD)

2017-09-13 06:27:25,350 [main] INFO org.apache.pig.Main - Pig script completed in 47 seconds and 647 milliseconds (47647 ms)

**gopalkrishna@ubuntu:~/PRAC/HBASE\$**

**ADD 1 MORE RECORD IN HBASE TABLE => AND RUN THE PIG SCRIPT AGAIN , WE WILL FIND THE NEW RECORD IN PIG SIDE**

hbase(main):002:0> put 'hbasepig\_emptable','1100','empdetails:firstname','RAJESH'  
0 row(s) in 0.0900 seconds

hbase(main):003:0> put 'hbasepig\_emptable','1100','empdetails:lastname','KANNA'  
0 row(s) in 0.0120 seconds

hbase(main):004:0> put 'hbasepig\_emptable','1100','empdetails:yrsofexp',14  
0 row(s) in 0.0080 seconds

hbase(main):005:0> put 'hbasepig\_emptable','1100','empdetails:desg','SOLUTION ARCHITECT'  
0 row(s) in 0.0130 seconds

hbase(main):006:0>

+++++

**FROM PIG SIDE – AFTER RUNNING THE SCRIPT AGAIN, WE WILL GET BELOW RESULTS**

Input(s):  
Successfully read 16 records (337 bytes) from: "hbase://hbasepig\_emptable"

Output(s):  
Successfully stored 16 records (777 bytes) in: "hdfs://localhost:8020/tmp/temp645148187/tmp-1399112497"

Counters:  
Total records written : 16  
Total bytes written : 777  
Spillable Memory Manager spill count : 0  
Total bags proactively spilled: 0  
Total records proactively spilled: 0

Job DAG:  
job\_1505297080421\_0008

2017-09-13 06:27:24,563 [main] INFO org.apache.hadoop.yarn.client.RMPProxy - Connecting to ResourceManager at localhost/127.0.0.1:8032  
2017-09-13 06:27:24,597 [main] INFO org.apache.hadoop.mapred.ClientServiceDelegate - Application state is completed. FinalApplicationStatus=SUCCEEDED. Redirecting to job history server  
2017-09-13 06:27:24,684 [main] INFO org.apache.hadoop.yarn.client.RMPProxy - Connecting to ResourceManager at localhost/127.0.0.1:8032  
2017-09-13 06:27:24,705 [main] INFO org.apache.hadoop.mapred.ClientServiceDelegate - Application state is completed. FinalApplicationStatus=SUCCEEDED. Redirecting to job history server  
2017-09-13 06:27:24,809 [main] INFO org.apache.hadoop.yarn.client.RMPProxy - Connecting to ResourceManager at localhost/127.0.0.1:8032  
2017-09-13 06:27:24,830 [main] INFO org.apache.hadoop.mapred.ClientServiceDelegate - Application state is completed. FinalApplicationStatus=SUCCEEDED. Redirecting to job history server  
2017-09-13 06:27:24,987 [main] INFO org.apache.pig.backend.hadoop.executionengine.mapReduceLayer.MapReduceLauncher - Success!  
2017-09-13 06:27:24,993 [main] INFO org.apache.hadoop.conf.Configuration.deprecation - fs.default.name is deprecated. Instead, use fs.defaultFS  
2017-09-13 06:27:25,001 [main] INFO org.apache.pig.data.SchemaTupleBackend - Key [pig.schematuple] was not set... will not generate code.

2017-09-13 06:27:25,095 [main] INFO org.apache.hadoop.mapreduce.lib.input.FileInputFormat - Total input paths to process : 1

2017-09-13 06:27:25,107 [main] INFO org.apache.pig.backend.hadoop.executionengine.util.MapRedUtil - Total input paths to process : 1

(1000,RAVI,VARMA,9,TECHNOLOGY LEAD)

(1001,KOUSIK,MISTRA,11,ENTRYLEVEL ARCHITECT)

(1002,SAIED,FAISAL,21,DELIVERY HEAD)

(1003,SOURAV,JAIN,18,SENIOR PROJECT MANAGER)

(1004,CHRIS,WAYNE,15,TEST MANAGER)

(1005,ANEES,FLEXWALA,19,CONFIGURATION CONTROLLER)

(1006,GOPAL,KRISHNA,13,SR.BIGDATA ARCHITECT)

(1007,RAMYA,SRI,8,BIGDATA LEAD)

(1008,MUNI,KRISHNA,5,SR.HADOOP DEVELOPER)

(1009,RAMYA,SRI,7,SR.SPARK DEVELOPER)

(1010,RAM,CHANDAR,16,PROJECT MANAGER)

(1011,AKHILA,PRIYA,3,HADOOP DEVELOPER)

(1012,RAHUL,DEV,6,HADOOP ADMINISTRATOR)

(1013,AMIT,MISRA,4,BIGDATA TESTER)

(1014,PUJA,YERUNKAR,11,BIGDATA TEST LEAD)

(1100,RAJESH,KANNA,14,SOLUTION ARCHITECT)

2017-09-13 06:27:25,350 [main] INFO org.apache.pig.Main - Pig script completed in 47 seconds and 647 milliseconds (47647 ms)

**gopalkrishna@ubuntu:~/PRAC/HBASE\$**